

# Digit\_Recog\_DNN

February 27, 2022

## 1 Final Scoring in the Competition and Ranking!

### Submission

YOUR RECENT SUBMISSION





**Michael\_Woo\_Predictions\_Digit\_Recog\_2.csv**

Submitted by Michael Woo · Submitted a day ago

Score: 0.99353

↓ Jump to your leaderboard position

### Leaderboard

230	Michael Woo		0.99353	2	1d
<div>Your Best Entry! Your most recent submission scored 0.99353, which is an improvement of your previous score of 0.72064. Great job!</div> <div>Tweet this</div>					
231	Gabriel Tavares		0.99350	5	9d
232	[Deleted] 4bf71481-4802-4c33-bbeb-6fbde14317bb		0.99350	23	6d

**Summary of Improvements** \* Since we are dealing with image data, I would think it would be best to use a few *convolutional* and *pooling* layers for this case \* Added a *Dropout layer*: To prevent overfitting \* Adjust *Neurons* per layer \* Added *Learning Rate Scheduling*: Piecewise Constant Scheduling \* Added *Batch Normalization*: To standardizes the inputs and stabilize the learning process \* Used activation function (*Elu*) and initialization of weights (*lecun\_normal*) \* Added more layers and neurons \* Used a different type of optimizer (*Adam*) \* Added a *gradient clipping*: To prevent exploding gradients

### Imports

```
[2]: import tensorflow as tf
from tensorflow import keras
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import Normalizer
```

## Dataset

```
[5]: digit = pd.read_csv("data/train.csv")
digit_test = pd.read_csv("data/test.csv")
```

## Training Dataset View

```
[6]: digit.head()
```

```
[6]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	\
0	1	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	
3	4	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	

	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	\
0	0	...	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	
2	0	...	0	0	0	0	0	0	
3	0	...	0	0	0	0	0	0	
4	0	...	0	0	0	0	0	0	

	pixel780	pixel781	pixel782	pixel783
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 785 columns]

## Test Dataset View

```
[7]: digit_test.head()
```

```
[7]:
```

	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	\
0	0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	
2	0	0	0	0	0	0	0	0	0	
3	0	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	

	pixel9	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	\
0	0	...	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	

2	0	...	0	0	0	0	0	0
3	0	...	0	0	0	0	0	0
4	0	...	0	0	0	0	0	0

	pixel780	pixel781	pixel782	pixel783
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 784 columns]

### Get only values

```
[8]: x = digit.drop('label', axis = 1).values
     y = digit.label.values
```

```
[9]: x.shape
```

```
[9]: (42000, 784)
```

```
[10]: y.shape
```

```
[10]: (42000,)
```

### Re-scaling

```
[11]: x = x/255.0
     test_df = digit_test.values/255.0
```

### Reshaping

```
[12]: x_train = x.reshape(x.shape[0], 28, 28,1)
     test_df = test_df.reshape(digit_test.shape[0], 28, 28,1)
     y_train = y.reshape(-1,1)
```

### Splitting Data

```
[13]: x_train, x_valid, y_train, y_valid = train_test_split(x_train,y_train,test_size=
     ↪ 0.20,
                                           random_state = 141)
```

### Clear the Backend

```
[18]: keras.backend.clear_session()
     np.random.seed(42)
     tf.random.set_seed(42)
```

**Piecewise Constant Scheduling** \* Added an additional learning constant \* A dynamic learning process

```
[19]: def piecewise_constant_fn(epoch):  
    if epoch < 5:  
        return 0.01  
    elif epoch < 10:  
        return 0.005  
    elif epoch < 15:  
        return 0.003  
    elif epoch < 20:  
        return 0.001  
    else:  
        return 0.0001
```

```
[20]: def piecewise_constant(boundaries, values):  
    boundaries = np.array([0] + boundaries)  
    values = np.array(values)  
    def piecewise_constant_fn(epoch):  
        return values[np.argmax(boundaries > epoch) - 1]  
    return piecewise_constant_fn  
piecewise_constant_fn = piecewise_constant([5,10,15,20,25], [0.01, 0.005,0.  
↪003,0.001,0.0001])
```

```
[21]: lr_scheduler = keras.callbacks.LearningRateScheduler(piecewise_constant_fn)
```

## Model Creation

```
[36]: model = keras.models.Sequential()  
model.add(keras.layers.Conv2D(filters=64, kernel_size=(5, 5),padding="valid" ,  
↪activation='elu',kernel_initializer="lecun_normal", input_shape=(28,28,1)))  
model.add(keras.layers.AveragePooling2D())  
model.add(keras.layers.BatchNormalization())  
model.add(keras.layers.Conv2D(filters=32, kernel_size=(5, 5),padding="valid",  
↪activation='elu',kernel_initializer="lecun_normal"))  
model.add(keras.layers.AveragePooling2D())  
model.add(keras.layers.BatchNormalization())  
model.add(keras.layers.GaussianDropout(0.25))  
model.add(keras.layers.Flatten())  
model.add(keras.layers.BatchNormalization())  
model.add(keras.layers.Dense(units=128,↪  
↪activation='selu',kernel_initializer="lecun_normal"))  
model.add(keras.layers.BatchNormalization())  
model.add(keras.layers.Dense(units=64,↪  
↪activation='selu',kernel_initializer="lecun_normal"))  
model.add(keras.layers.Dense(units=10, activation = 'softmax'))
```

```
[37]: model.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 24, 24, 64)	1664
average_pooling2d_4 (Average Pooling2D)	(None, 12, 12, 64)	0
batch_normalization_14 (Batch Normalization)	(None, 12, 12, 64)	256
conv2d_7 (Conv2D)	(None, 8, 8, 32)	51232
average_pooling2d_5 (Average Pooling2D)	(None, 4, 4, 32)	0
batch_normalization_15 (Batch Normalization)	(None, 4, 4, 32)	128
gaussian_dropout_3 (Gaussian Dropout)	(None, 4, 4, 32)	0
flatten_3 (Flatten)	(None, 512)	0
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dense_13 (Dense)	(None, 128)	65664
batch_normalization_17 (Batch Normalization)	(None, 128)	512
dense_14 (Dense)	(None, 64)	8256
dense_15 (Dense)	(None, 10)	650

```
=====  
Total params: 130,410  
Trainable params: 128,938  
Non-trainable params: 1,472  
=====
```

```
-----  
Model Compile
```

```
[38]: model.compile(loss="sparse_categorical_crossentropy",optimizer=keras.optimizers.  
      ↪Adam(clipvalue=1.0,clipnorm=1.0),metrics=["accuracy"])
```

## Model Fitting

```
[39]: history = model.  
      ↪fit(x_train,y_train,epochs=25,validation_data=(x_valid,y_valid),callbacks=[lr_scheduler])
```

Epoch 1/25

```
2022-02-26 15:25:47.913170: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
1049/1050 [=====>.] - ETA: 0s - loss: 0.2338 - accuracy:  
0.9304
```

```
2022-02-26 15:26:21.398391: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.
```

```
1050/1050 [=====] - 37s 34ms/step - loss: 0.2337 -  
accuracy: 0.9304 - val_loss: 0.0989 - val_accuracy: 0.9731 - lr: 0.0100
```

Epoch 2/25

```
1050/1050 [=====] - 33s 32ms/step - loss: 0.1190 -  
accuracy: 0.9651 - val_loss: 0.0793 - val_accuracy: 0.9771 - lr: 0.0100
```

Epoch 3/25

```
1050/1050 [=====] - 33s 32ms/step - loss: 0.1025 -  
accuracy: 0.9709 - val_loss: 0.2088 - val_accuracy: 0.9464 - lr: 0.0100
```

Epoch 4/25

```
1050/1050 [=====] - 34s 33ms/step - loss: 0.0933 -  
accuracy: 0.9731 - val_loss: 0.0791 - val_accuracy: 0.9789 - lr: 0.0100
```

Epoch 5/25

```
1050/1050 [=====] - 33s 32ms/step - loss: 0.0853 -  
accuracy: 0.9770 - val_loss: 0.0539 - val_accuracy: 0.9880 - lr: 0.0100
```

Epoch 6/25

```
1050/1050 [=====] - 34s 32ms/step - loss: 0.0467 -  
accuracy: 0.9856 - val_loss: 0.0412 - val_accuracy: 0.9890 - lr: 0.0050
```

Epoch 7/25

```
1050/1050 [=====] - 33s 32ms/step - loss: 0.0426 -  
accuracy: 0.9866 - val_loss: 0.0419 - val_accuracy: 0.9880 - lr: 0.0050
```

Epoch 8/25

```
1050/1050 [=====] - 33s 31ms/step - loss: 0.0438 -  
accuracy: 0.9861 - val_loss: 0.0466 - val_accuracy: 0.9864 - lr: 0.0050
```

Epoch 9/25

```
1050/1050 [=====] - 33s 32ms/step - loss: 0.0375 -  
accuracy: 0.9884 - val_loss: 0.0426 - val_accuracy: 0.9888 - lr: 0.0050
```

Epoch 10/25

```
1050/1050 [=====] - 34s 32ms/step - loss: 0.0390 -  
accuracy: 0.9884 - val_loss: 0.0393 - val_accuracy: 0.9905 - lr: 0.0050
```

Epoch 11/25  
1050/1050 [=====] - 34s 32ms/step - loss: 0.0270 -  
accuracy: 0.9917 - val\_loss: 0.0301 - val\_accuracy: 0.9917 - lr: 0.0030

Epoch 12/25  
1050/1050 [=====] - 34s 32ms/step - loss: 0.0230 -  
accuracy: 0.9930 - val\_loss: 0.0342 - val\_accuracy: 0.9918 - lr: 0.0030

Epoch 13/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0235 -  
accuracy: 0.9927 - val\_loss: 0.0342 - val\_accuracy: 0.9921 - lr: 0.0030

Epoch 14/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0226 -  
accuracy: 0.9927 - val\_loss: 0.0363 - val\_accuracy: 0.9911 - lr: 0.0030

Epoch 15/25  
1050/1050 [=====] - 419s 399ms/step - loss: 0.0199 -  
accuracy: 0.9932 - val\_loss: 0.0356 - val\_accuracy: 0.9901 - lr: 0.0030

Epoch 16/25  
1050/1050 [=====] - 34s 32ms/step - loss: 0.0147 -  
accuracy: 0.9953 - val\_loss: 0.0307 - val\_accuracy: 0.9918 - lr: 0.0010

Epoch 17/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0123 -  
accuracy: 0.9958 - val\_loss: 0.0342 - val\_accuracy: 0.9906 - lr: 0.0010

Epoch 18/25  
1050/1050 [=====] - 34s 32ms/step - loss: 0.0107 -  
accuracy: 0.9963 - val\_loss: 0.0331 - val\_accuracy: 0.9915 - lr: 0.0010

Epoch 19/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0103 -  
accuracy: 0.9965 - val\_loss: 0.0335 - val\_accuracy: 0.9921 - lr: 0.0010

Epoch 20/25  
1050/1050 [=====] - 55s 52ms/step - loss: 0.0103 -  
accuracy: 0.9968 - val\_loss: 0.0330 - val\_accuracy: 0.9929 - lr: 0.0010

Epoch 21/25  
1050/1050 [=====] - 120s 114ms/step - loss: 0.0080 -  
accuracy: 0.9975 - val\_loss: 0.0306 - val\_accuracy: 0.9925 - lr: 1.0000e-04

Epoch 22/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0087 -  
accuracy: 0.9970 - val\_loss: 0.0303 - val\_accuracy: 0.9924 - lr: 1.0000e-04

Epoch 23/25  
1050/1050 [=====] - 34s 32ms/step - loss: 0.0078 -  
accuracy: 0.9973 - val\_loss: 0.0302 - val\_accuracy: 0.9927 - lr: 1.0000e-04

Epoch 24/25  
1050/1050 [=====] - 33s 32ms/step - loss: 0.0078 -  
accuracy: 0.9971 - val\_loss: 0.0295 - val\_accuracy: 0.9926 - lr: 1.0000e-04

Epoch 25/25  
1050/1050 [=====] - 123s 117ms/step - loss: 0.0078 -  
accuracy: 0.9974 - val\_loss: 0.0285 - val\_accuracy: 0.9924 - lr: 1.0000e-04

**Save the Model** \* I already saved this model to local and cloud

```
[112]: #model.save("Digit_Recog_Mod_2.h5")
```

```
[3]: ml = keras.models.load_model("Digit_Recog_Mod_2.h5")
```

```
2022-02-27 20:54:07.840888: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
```

```
2022-02-27 20:54:07.841283: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)
```

Metal device set to: Apple M1 Pro

### Testing Model Against Test Data

```
[14]: pred = ml.predict(test_df)
```

```
2022-02-27 20:54:44.131388: W
tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU
frequency: 0 Hz
```

```
2022-02-27 20:54:44.193648: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

Get the index of the max value in the array

```
[15]: prediction = [np.argmax(i) for i in pred]
```

```
[16]: img_id = list(range(1,28001))
img_id = np.array(img_id)
```

```
[17]: img_id.shape
```

```
[17]: (28000,)
```

**Final Dataframe of Predictions** \* Has to be in this format for the competition

```
[18]: df = pd.DataFrame({"ImageId":img_id,"Label":prediction})
df.head()
```

```
[18]:   ImageId  Label
0         1      2
1         2      0
2         3      9
3         4      0
4         5      3
```



```
[19]: df.head()
```

```
[19]:
```

	ImageId	Label
0	1	2
1	2	0
2	3	9
3	4	0
4	5	3

```
[20]: df.tail()
```

```
[20]:
```

	ImageId	Label
27995	27996	9
27996	27997	7
27997	27998	3
27998	27999	9
27999	28000	2

**Download to csv format**

```
[108]: df.to_csv("Michael_Woo_Predictions_Digit_Recog_2.csv", index=False)
```